

Outline of the notebook

In this notebook, I will show my codes and workings for Natural Language Processing (NLP) techniques used to study the Airbnb reviews. The dataset contains more than 100,000 reviews left by guests who had stayed in Singapore listings in the past.

The models used include:

- Bag of Words
- Sentiment Analysis
- Topic Modelling

The libraries used include:

- LangDetect
- CountVectorizer
- Textblob
- Vader
- Gensim

Import basic packages that will be used

```
In [ ]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt

%matplotlib inline
plt.style.use('fivethirtyeight')
%config InlineBackend.figure_format = 'retina'
```

Load the datasets to be used

```
In [2]: sglisting = pd.read_csv('./SG listings (1).csv')
```

```
In [2]: sgreviews = pd.read_csv('./SG reviews.csv')
```

```
In [291]: sgreviews.shape
```

```
Out[291]: (101009, 6)
```

```
In [66]: sgreviews.head()
```

```
Out[66]:
```

	listing_id	id	date	reviewer_id	reviewer_name	comments	all_cleaned
0	49091	8243238	2013-10-21	8557223	Jared	Fran was absolutely gracious and welcoming. Ma...	Fran was absolutely gracious and welcoming. Ma... Franwasabsce
1	50646	11909864	2014-04-18	1356099	James	A comfortable room in a smart condo developmen...	A comfortable room in a smart condo developmen... Acomfortabl
2	50646	13823948	2014-06-05	15222393	Welli	Stayed over at Sujatha's house for 3 good nigh...	Stayed over at Sujatha's house for 3 good nigh... Stayedover
3	50646	15117222	2014-07-02	5543172	Cyril	It's been a lovely stay at Sujatha's. The room...	It's been a lovely stay at Sujatha's. The room... Itsbeena
4	50646	15426462	2014-07-08	817532	Jake	We had a great experience. A nice place, an am...	We had a great experience. A nice place, an am... Wehadagrea


Visualizing and studying review scores

Review ratings are provided in the listings dataset.

```
In [80]: # Plot overall reviews scores

sns.distplot(sglisting['review_scores_rating'].dropna())
```


```
Out[80]: <matplotlib.axes._subplots.AxesSubplot at 0x1a1f5ad6d8>
```



```
In [81]: # Plot reviews scores
```

```
sns.distplot(sglisting['review_scores_accuracy'].dropna())
```


```
Out[81]: <matplotlib.axes._subplots.AxesSubplot at 0x1a1f5ec940>
```



```
In [82]: # Plot overall reviews scores
```

```
sns.distplot(sglisting['review_scores_cleanliness'].dropna())
```


```
Out[82]: <matplotlib.axes._subplots.AxesSubplot at 0x1a1f62d978>
```



```
In [83]: # Plot overall reviews scores
```

```
sns.distplot(sglisting['review_scores_checkin'].dropna())
```


```
Out[83]: <matplotlib.axes._subplots.AxesSubplot at 0x1a2017f438>
```



```
In [84]: # Plot overall reviews scores
```

```
sns.distplot(sglisting['review_scores_communication'].dropna())
```


```
Out[84]: <matplotlib.axes._subplots.AxesSubplot at 0x1a18ab2518>
```



```
In [85]: # Plot overall reviews scores
```

```
sns.distplot(sglisting['review_scores_location'].dropna())
```


```
Out[85]: <matplotlib.axes._subplots.AxesSubplot at 0x1a20b86860>
```



```
In [86]: # Plot overall reviews scores
```

```
sns.distplot(sglisting['review_scores_value'].dropna())
```

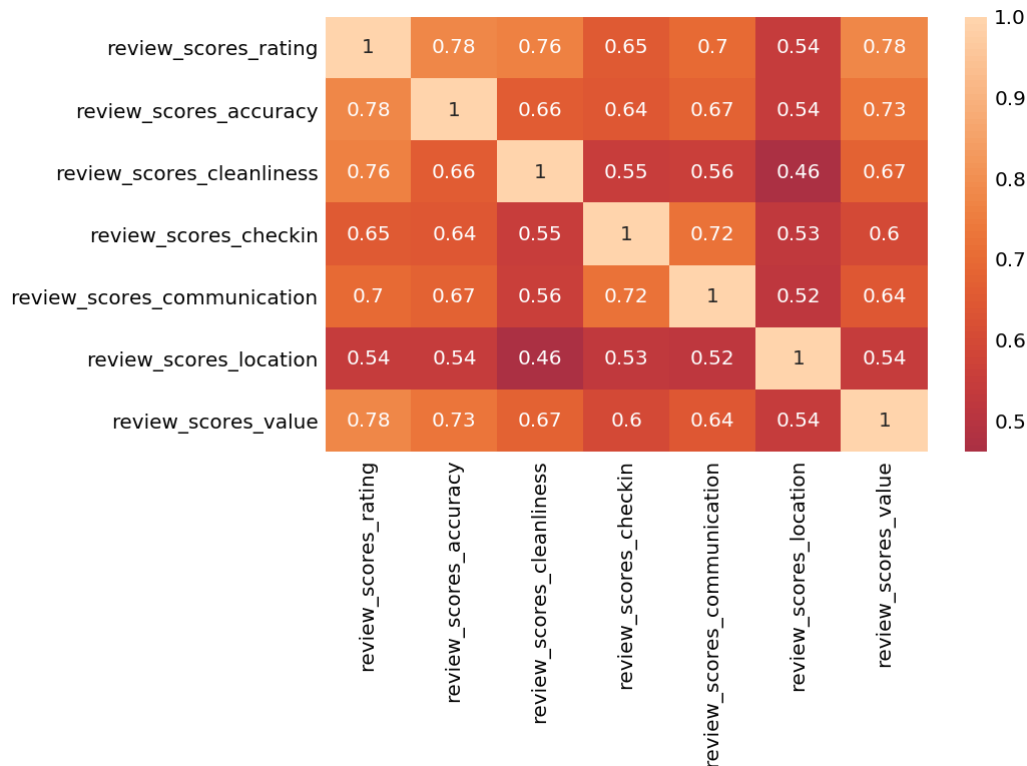
```
Out[86]: <matplotlib.axes._subplots.AxesSubplot at 0x1a2105abe0>
```



Guests tend to rate value lower than the other portions - customers always demand more value (more benefits for a lower price).

```
In [90]: # Plot correlation between the different scores

sns.heatmap(sglisting[['review_scores_rating', 'review_scores_accuracy', 'review_scores_cleanliness', 'review_scores_checkin', 'review_scores_communication', 'review_scores_location', 'review_scores_value']].corr(), center=0, annot=True)
```



```
In [91]: # Find the average score for different features

sglisting[['review_scores_rating', 'review_scores_accuracy', 'review_scores_cleanliness', 'review_scores_checkin', 'review_scores_communication', 'review_scores_location', 'review_scores_value']].mean()
```

```
Out[91]: review_scores_rating      90.700226
review_scores_accuracy      9.252925
review_scores_cleanliness    9.033846
review_scores_checkin        9.523017
review_scores_communication   9.496202
review_scores_location        9.324152
review_scores_value          9.021167
dtype: float64
```

Data Cleaning - Reviews dataset

Drop null reviews

We are unable to obtain any information from the null reviews in this case

```
In [3]: # Drop all null values in sgreviews

sgreviews.dropna(inplace = True)
```

Drop all automated reviews

```
In [4]: # Find common text present for automated reviews due to cancellation
# Filter them out to remove them in the next step

txt = 'This is an automated posting'

# Get list of actual reviews vs automated reviews
cleaned_review = []
automated_review = []

for value in sgreviews['comments'].values:
    if txt in value:
        automated_review.append(value)
    else:
        cleaned_review.append(value)

# Get list of all reviews, automated reviews as null values
all_reviews_cleaned = []

for value in sgreviews['comments'].values:
    if txt not in value:
        all_reviews_cleaned.append(value)
    else:
        all_reviews_cleaned.append(np.nan)
```

```
In [5]: # Count number of automated reviews vs all reviews

print('{} {} {}'.format('Number of automated reviews: ', len(automated_review), len(cleaned_review)))
print('{} {} {}'.format('Number of actual reviews: ', len(cleaned_review), len(automated_review)))
print('{} {} {}'.format('Number of total reviews: ', len(automated_review) + len(cleaned_review)))

Number of automated reviews: 1530 1.5159471697365423%
Number of actual reviews: 99397 98.48405283026345%
Number of total reviews: 100927
```

```
In [6]: # Create new column for cleaned reviews

sgreviews['all_cleaned'] = all_reviews_cleaned
```

```
In [7]: # Drop automated reviews

sgreviews.dropna(inplace=True)
```

Check for reviews in foreign languages & Extract only English reviews

```
In [8]: # Import the library for detecting languages
```

```
import langdetect
```

```
In [9]: # List of languages under langdetect
```

```
lang = ['af', 'ar', 'bg', 'bn', 'ca', 'cs', 'cy', 'da', 'de', 'el', 'en', 'gu', 'he', 'hi', 'hr', 'hu', 'id', 'it', 'ja', 'kn', 'ko', 'lt', 'nl', 'no', 'pa', 'pl', 'pt', 'ro', 'ru', 'sk', 'sl', 'so', 'sq', 'tl', 'tr', 'uk', 'ur', 'vi', 'zh']
```

```
In [10]: cleaned_list = []
```

```
for review in sgreviews['all_cleaned'].values:  
    cleaned_list.append(review.strip())
```

```
In [11]: import re
```

```
cleaner_list = []
```

```
for review in cleaned_list:  
    cleaner_list.append(re.sub("[ !@#$%^&*()_+==.,<>?:;|}{?~`'/1234567890]"
```

```
sgreviews['cleaner_list'] = cleaner_list
```

```
drop_list = sgreviews[sgreviews['cleaner_list'] == ''].index  
sgreviews = sgreviews.drop(drop_list, axis=0)
```

```
In [12]: # Retrieve languages for the reviews using langdetect
```

```
lang_list = []
```

```
for i,value in enumerate(sgreviews['all_cleaned'].values):  
    try:  
        lang_list.append(langdetect.detect_langs(value))  
    except:  
        lang_list.append(langdetect.detect_langs('Undetectable'))
```

```
In [13]: # Extract the most probable language
```

```
cleaned_lang_list = []
```

```
for i, value in enumerate(lang_list):  
    if value == 'Undetectable':  
        cleaned_lang_list.append(np.nan)  
    else:  
        cleaned_lang_list.append(str(lang_list[i][0])[2])
```

```
In [14]: # Add list of languages back into sgreviews dataframe
```

```
sgreviews['review_lang'] = cleaned_lang_list
```



```
In [15]: # Assign new variable en_reviews for reviews that are in english

en_reviews = sgreviews[sgreviews['review_lang'] == 'en']
```

Sentiment Analysis on reviews

```
In [16]: # Create function for sentiment analysis using TextBlob

from textblob import TextBlob, Word

def sentiment_analysis_blob(feature):

    sentiment_blob = []

    # Score text using TextBlob sentiment polarity attribute
    for listing in en_reviews[feature].values:
        try:
            sentiment_blob.append(TextBlob(listing.lower()).sentiment.polarity)
        except AttributeError:
            sentiment_blob.append(0)

    return sentiment_blob
```

```
In [17]: # Create function for sentiment analysis using Vader

from vaderSentiment.vaderSentiment import SentimentIntensityAnalyzer

def sentiment_analysis_vader(feature):

    # Score text using Vader SentimentIntensityAnalyzer
    analyzer = SentimentIntensityAnalyzer()
    sentiment_vader = []

    for listing in en_reviews[feature].values:
        try:
            sentiment_vader.append(analyzer.polarity_scores(listing.lower())['compound'])
        except AttributeError:
            sentiment_vader.append({'compound': 0.0})

    sentiment_vader_compound = []

    for listing in sentiment_vader:
        sentiment_vader_compound.append(listing['compound'])

    return sentiment_vader_compound
```

```
In [18]: # Use the function on features
```

```
en_reviews['comments_blob'] = sentiment_analysis_blob('all_cleaned')
en_reviews['comments_vader'] = sentiment_analysis_vader('all_cleaned')
```

/anaconda3/lib/python3.6/site-packages/ipykernel_launcher.py:3: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy> (<http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>)

This is separate from the ipykernel package so we can avoid doing imports until

/anaconda3/lib/python3.6/site-packages/ipykernel_launcher.py:4: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy> (<http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>)

after removing the cwd from sys.path.

```
In [19]: # Average score from blob and vader
```

```
en_reviews['comments_ave'] = (en_reviews['comments_blob'] + en_reviews['com
```

/anaconda3/lib/python3.6/site-packages/ipykernel_launcher.py:3: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy> (<http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>)

This is separate from the ipykernel package so we can avoid doing imports until

Plot distribution of polarity among ENGLISH reviews


```
In [183]: # Plot distribution of sentiment polarity
```

```
sns.distplot(en_reviews['comments_blob'])
```

/anaconda3/lib/python3.6/site-packages/scipy/stats/stats.py:1713: FutureWarning: Using a non-tuple sequence for multidimensional indexing is deprecated; use `arr[tuple(seq)]` instead of `arr[seq]`. In the future this will be interpreted as an array index, `arr[np.array(seq)]`, which will result either in an error or a different result.

```
return np.add.reduce(sorted[indexer] * weights, axis=axis) / sumval
```

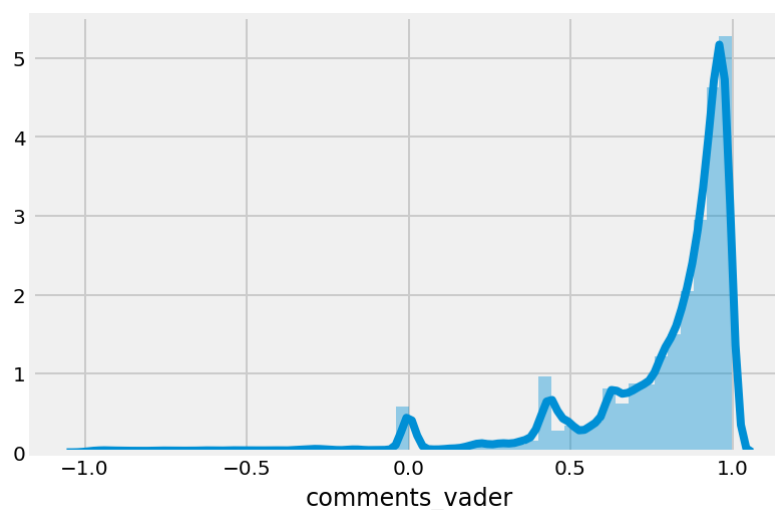
```
Out[183]: <matplotlib.axes._subplots.AxesSubplot at 0x1a1be084e0>
```



```
In [185]: # Plot distribution of sentiment polarity
```

```
sns.distplot(en_reviews['comments_vader'])
```


```
Out[185]: <matplotlib.axes._subplots.AxesSubplot at 0x1a2c486cc0>
```



```
In [65]: # Plot distribution of sentiment polarity

sns.distplot(en_reviews['comments_ave'])
plt.title('Distribution of Sentiment Polarity Scores')
plt.ylabel('Frequency')
plt.xlabel('Polarity Score')
```

```
Out[65]: Text(0.5, 0, 'Polarity Score')
```



```
In [299]: # Check the reviews text for the most negative reviews

en_reviews[['all_cleaned', 'comments_ave']].sort_values(by='comments_ave', a
```

```
Out[299]:
```

	all_cleaned	comments_ave
81355	짬쥬이니 저한테 실슈를 해노코 요히료 엄마를 차자가게따 공찰을 뷰루겠다 흠뻑했여오....	-0.926850
91102	It was the worst airbnb experience that I've e...	-0.812450
77437	Worst stay ever	-0.812450
57902	worst experience ever	-0.812450
71357	The worst ever people trying to con people's m...	-0.812450
37694	Very bad! Whill I asking about canceling my bo...	-0.811500
78813	Location is awful, takes quite a while to get ...	-0.775300
19617	Horrible. Will not going there again	-0.771150
74295	Such a bad room. The room was so dirty and inc...	-0.762283
99211	The place is not as advertised. I had lots of ...	-0.754783
30102	Worst and Terrible Host. Not for recommendatio...	-0.751333
15050	poor communication and worst reception.	-0.751000
54149	Wifi very bad	-0.747450
86654	This guy canceled on us WHILE we were looking ...	-0.747450
45086	it is a horrible horrible experience that the ...	-0.744100

Typesetting math: 0%

```
In [46]: # Average polarity score for the English reviews

en_reviews['comments_ave'].mean()
```

```
Out[46]: 0.5743625682821794
```

Bag of Words on all reviews

```
In [23]: # Load NLP libraries

from sklearn.feature_extraction.text import CountVectorizer
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.feature_extraction import stop_words
import nltk
from nltk.corpus import stopwords
from nltk.tokenize import word_tokenize, sent_tokenize
from textblob import TextBlob, Word
from wordcloud import WordCloud, STOPWORDS, ImageColorGenerator
```

```
In [24]: # Create function to get frequent words for CountVectorizer / TF-IDF Vectorizer

def get_freq_words(sparse_counts, columns):
    # X_all is a sparse matrix, so sum() returns a 'matrix' datatype ...
    #   which we then convert into a 1-D ndarray for sorting
    word_counts = np.asarray(sparse_counts.sum(axis=0)).reshape(-1)

    # argsort() returns smallest first, so we reverse the result
    largest_count_indices = word_counts.argsort()[::-1]

    # pretty-print the results! Remember to always ask whether they make sense
    freq_words = pd.Series(word_counts[largest_count_indices],
                           index=columns[largest_count_indices])

    return freq_words
```

```
In [27]: Retrieve default stopwords in the NLTK library
```

```
stop = stopwords.words('english')
```

```
Add more stopwords that are applicable in this case
```

```
stop.extend(['stay', 'recommend', 'recommended', 'nice', 'definitely', 'great', 'please'])
```

```
In [31]: cvec = CountVectorizer(stop_words=stop, ngram_range=(2,4), min_df=80, max_df=10)
cvec.fit(en_reviews['all_cleaned'])
df_train = pd.DataFrame(cvec.transform(en_reviews['all_cleaned']).todense(),
                        columns=cvec.get_feature_names())
df_test = pd.DataFrame(cvec.transform(en_reviews['all_cleaned']).todense(),
                      columns=cvec.get_feature_names())
columns = np.array(cvec.get_feature_names())
```

```
In [32]: get_freq_words(cvec.transform(en_reviews['all_cleaned']), columns)
```

```
Out[32]: mrt station 5571
walking distance 3528
bus stop 2731
room clean 2515
close mrt 2103
value money 1936
clean comfortable 1798
public transport 1768
near mrt 1760
minutes walk 1737
apartment clean 1724
swimming pool 1666
location close 1524
everything need 1480
little india 1477
easy access 1394
minute walk 1372
walk mrt 1366
friendly helpful 1306
easy get 1277
orchard road 1244
visit singapore 1190
convenient location 1190
mins walk 1173
washing machine 1160
location near 1139
gave us 1134
min walk 1126
bus stops 1080
clean well 1070
...
singapore near 81
see pictures 81
friendly always 81
strategically located 81
like local 81
location easy find 81
thank hosting 81
working well 81
location wonderful 81
places eat nearby 80
helpful information 80
singapore family 80
solo traveller 80
whole apartment 80
house spacious 80
pool roof 80
convenient transportation 80
stop right outside 80
little noisy 80
clean check 80
clean facilities 80
mrt restaurants 80
room basic 80
fast respond 80
```

Typesetting math: 99%

```

room clean well      80
clean lovely         80
room convenient      80
make sure everything 80
heavy luggage        80
get see              80
Length: 1866, dtype: int64

```

Bag of Words on Positive Reviews

In [134]: Retrieve default stopwords in the NLTK library

```
stop = stopwords.words('english')
```

Add more stopwords that are applicable in this case

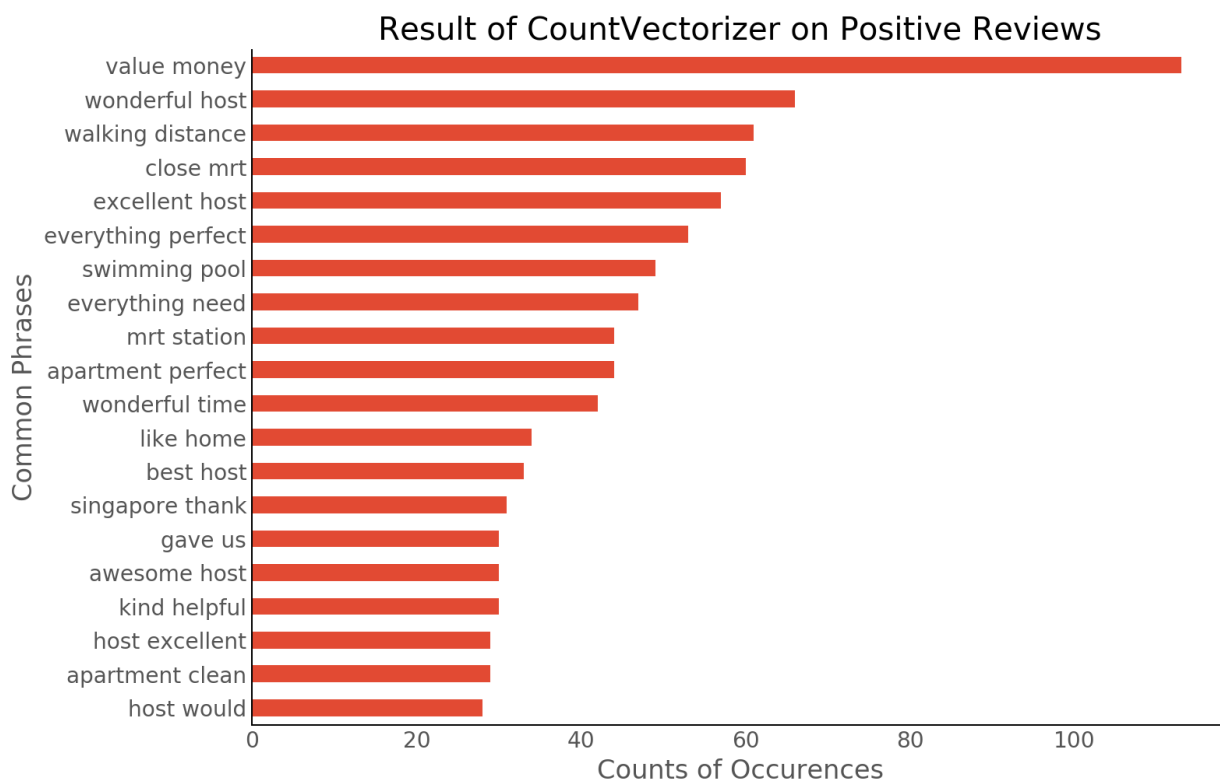
```
stop.extend(['stay', 'recommend', 'recommended', 'nice', 'definitely', 'great', 'pl
```

```

In [135]: cvec = CountVectorizer(stop_words=stop, ngram_range=(2,4), min_df=10, max_d
cvec.fit(en_reviews[en_reviews['comments_ave'] > 0.8]['all_cleaned'])
df_train = pd.DataFrame(cvec.transform(en_reviews[en_reviews['comments_ave
columns=cvec.get_feature_names()
df_test = pd.DataFrame(cvec.transform(en_reviews[en_reviews['comments_ave']
columns=cvec.get_feature_names()
columns = np.array(cvec.get_feature_names())

```

```
In [136]: plt.rcParams.update({'font.size': 14})
fig, ax = plt.subplots(1,1)
get_freq_words(cvec.transform(en_reviews[en_reviews['comments_ave'] > 0.8])
plt.xlabel('Counts of Occurences')
plt.ylabel('Common Phrases')
plt.title('Result of CountVectorizer on Positive Reviews')
ax.spines['bottom'].set_color('black')
ax.spines['left'].set_color('black')
fig.set_facecolor('white')
ax.set_facecolor('white');
```




```
In [104]: # Retrieve default stopwords in the NLTK library

stop = stopwords.words('english')

# Add more stopwords that are applicable in this case
stop.extend(['stay', 'recommend', 'recommended', 'nice', 'definitely', 'great', 'website', 'hidden', 'airbnb', 'mrt', 'bus', 'mins', 'bnb', 'location', 'experience', 'first'])
```

```
In [105]: cvec = CountVectorizer(stop_words=stop, ngram_range=(2,4), min_df=10, max_df=10)
cvec.fit(en_reviews[en_reviews['comments_ave'] < 0]['all_cleaned'])
df_train = pd.DataFrame(cvec.transform(en_reviews[en_reviews['comments_ave'] < 0]['all_cleaned']).toarray(),
                        columns=cvec.get_feature_names())
df_test = pd.DataFrame(cvec.transform(en_reviews[en_reviews['comments_ave'] > 0]['all_cleaned']).toarray(),
                      columns=cvec.get_feature_names())
columns = np.array(cvec.get_feature_names())
```

```
In [106]: get_freq_words(cvec.transform(en_reviews[en_reviews['comments_ave'] < 0])['a
```

```
Out[106]: living room          78
washing machine        61
hot water              56
room small             46
toilet paper           46
air conditioning       45
little india           45
room clean             37
check time             33
one night              31
little bit             30
sofa bed               30
small room             29
told us                27
late night             25
near stop              25
value money            24
open door              24
late check             24
common area            23
air conditioner        23
clean room             22
one room               22
last minute            21
really bad             20
room bathroom          20
gave us                20
last day               20
one bathroom           19
every time             19
..
quite noisy            11
overall bad            11
one thing              11
one bed                11
never met              11
nearest station        11
one person             11
find another           11
however room           11
farrer park            11
check easy             10
air condition          10
working properly       10
day check              10
washer dryer           10
extremely small        10
arrived late           10
feel like              10
basic amenities        10
bathroom toilet        10
look like              10
size bed               10
room kitchen           10
get refund             10
```

Typesetting math: 9%

```


host said      10
close stop     10
different room 10
living area    10
long time      10
worth money    10
Length: 157, dtype: int64

```

```

In [128]: plt.rcParams.update({'font.size': 14})
fig, ax = plt.subplots(1,1)
plt.style.use('ggplot')
get_freq_words(cvec.transform(en_reviews[en_reviews['comments_ave'] < 0]['a
plt.xlabel('Counts of Occurences')
plt.ylabel('Common Phrases')
plt.title('Result of CountVectorizer on Negative Reviews')
ax.spines['bottom'].set_color('black')
ax.spines['left'].set_color('black')
fig.set_facecolor('white')
ax.set_facecolor('white');

```



Topic Modelling on Reviews

```

In [20]: # Import gensim package

import gensim
import gensim.corpora as corpora
from gensim.utils import simple_preprocess
from gensim.models import CoherenceModel

```

```
In [21]: # spacy for lemmatization
import spacy
```

```
In [33]: # Convert to list
data = en_reviews['all_cleaned'].values.tolist()
```

```
In [34]: # Tokenize all the reviews (remove punctuations, make lowercase)

def sent_to_words(sentences):
    for sentence in sentences:
        yield(gensim.utils.simple_preprocess(str(sentence), deacc=True)) #

data_words = list(sent_to_words(data))

# Check tokenized reviews
print(data_words[:1])

[['fran', 'was', 'absolutely', 'gracious', 'and', 'welcoming', 'made', 'm
y', 'stay', 'great', 'experience', 'would', 'definitely', 'recommend', 't
his', 'cozy', 'and', 'peaceful', 'place', 'to', 'anyone']]
```

```
In [35]: # Build the bigram and trigram models
bigram = gensim.models.Phrases(data_words, min_count=5, threshold=50) # hig
trigram = gensim.models.Phrases(bigram[data_words], threshold=50)

# Faster way to get a sentence clubbed as a trigram/bigram
bigram_mod = gensim.models.phrases.Phraaser(bigram)
trigram_mod = gensim.models.phrases.Phraaser(trigram)

# See trigram example
print(trigram_mod[bigram_mod[data_words[0]]])
```

```
/anaconda3/lib/python3.6/site-packages/gensim/models/phrases.py:494: User
Warning: For a faster implementation, use the gensim.models.phrases.Phras
er class
  warnings.warn("For a faster implementation, use the gensim.models.phras
es.Phraaser class")
/anaconda3/lib/python3.6/site-packages/gensim/models/phrases.py:494: User
Warning: For a faster implementation, use the gensim.models.phrases.Phras
er class
  warnings.warn("For a faster implementation, use the gensim.models.phras
es.Phraaser class")
/anaconda3/lib/python3.6/site-packages/gensim/models/phrases.py:494: User
Warning: For a faster implementation, use the gensim.models.phrases.Phras
er class
  warnings.warn("For a faster implementation, use the gensim.models.phras
es.Phraaser class")
/anaconda3/lib/python3.6/site-packages/gensim/models/phrases.py:494: User
Warning: For a faster implementation, use the gensim.models.phrases.Phras
er class
  warnings.warn("For a faster implementation, use the gensim.models.phras
es.Phraaser class")
```

```
In [36]: # Define functions for stopwords, bigrams, trigrams and lemmatization
def remove_stopword(texts):
    return [[word for word in simple_preprocess(str(doc)) if word not in st

def make_bigrams(texts):
    return [bigram_mod[doc] for doc in texts]

def make_trigrams(texts):
    return [trigram_mod[bigram_mod[doc]] for doc in texts]

def lemmatization(texts, allowed_postags=['NOUN', 'ADJ', 'VERB', 'ADV']):
    """https://spacy.io/api/annotation"""
    texts_out = []
    for sent in texts:
        doc = nlp(" ".join(sent))
        texts_out.append([token.lemma_ for token in doc if token.pos_ in al
    return texts_out
```

```
In [37]: # Remove Stop Words
data_words_nostops = remove_stopword(data_words)

# Form Bigrams
data_words_bigrams = make_bigrams(data_words_nostops)

# Initialize spacy 'en' model, keeping only tagger component (for efficiency)
# python3 -m spacy download en
nlp = spacy.load('en', disable=['parser', 'ner'])

# Do lemmatization keeping only noun, adj, vb, adv
data_lemmatized = lemmatization(data_words_bigrams, allowed_postags=['NOUN'

print(data_lemmatized[:1])

[['fran', 'absolutely', 'gracious', 'welcome', 'make', 'experience', 'wou
ld', 'cozy', 'peaceful', 'anyone']]
```

```
In [38]: # Create Dictionary
id2word = corpora.Dictionary(data_lemmatized)

# Create Corpus
texts = data_lemmatized

# Term Document Frequency
corpus = [id2word.doc2bow(text) for text in texts]

# View
print(corpus[:1])

[[ (0, 1), (1, 1), (2, 1), (3, 1), (4, 1), (5, 1), (6, 1), (7, 1), (8, 1),
(9, 1) ]]
```

```
In [254]: # Human readable format of corpus (term-frequency)
[[id2word[id], freq] for id, freq in cp] for cp in corpus[:1]]
```

```
Out[254]: [('absolutely', 1),
 ('anyone', 1),
 ('cozy', 1),
 ('experience', 1),
 ('fran', 1),
 ('gracious', 1),
 ('make', 1),
 ('peaceful', 1),
 ('welcome', 1),
 ('would', 1)]
```

```
In [271]: # Build LDA model - try with 10 topics
lda_model = gensim.models.ldamodel.LdaModel(corpus=corpus,
                                             id2word=id2word,
                                             num_topics=10,
                                             random_state=100,
                                             update_every=1,
                                             chunksize=100,
                                             passes=10,
                                             alpha='auto',
                                             per_word_topics=True)
```

recreationWarning: Calling np.sum(generator) is deprecated, and in the future will give a different result. Use np.sum(np.from_iter(generator)) or the python sum builtin instead.

score += np.sum(cnt * logsumexp(Elogtheta + Elogbeta[:, int(id)])) for id, cnt in doc)

/anaconda3/lib/python3.6/site-packages/gensim/models/ldamodel.py:826: DeprecationWarning: Calling np.sum(generator) is deprecated, and in the future will give a different result. Use np.sum(np.from_iter(generator)) or the python sum builtin instead.

score += np.sum(cnt * logsumexp(Elogtheta + Elogbeta[:, int(id)])) for id, cnt in doc)

/anaconda3/lib/python3.6/site-packages/gensim/models/ldamodel.py:826: DeprecationWarning: Calling np.sum(generator) is deprecated, and in the future will give a different result. Use np.sum(np.from_iter(generator)) or the python sum builtin instead.

score += np.sum(cnt * logsumexp(Elogtheta + Elogbeta[:, int(id)])) for id, cnt in doc)

/anaconda3/lib/python3.6/site-packages/gensim/models/ldamodel.py:826: DeprecationWarning: Calling np.sum(generator) is deprecated, and in the future will give a different result. Use np.sum(np.from_iter(generator)) or the python sum builtin instead.

```
In [272]: # Print the Keyword in the 10 topics
```

```
from pprint import pprint
```

```
pprint(lda_model.print_topics())
doc_lda = lda_model[corpus]
```

```
[(0,
  '0.099*"check" + 0.047*"night" + 0.045*"super" + 0.041*"even" + '
  '0.034*"excellent" + 0.033*"pool" + 0.029*"lovely" + 0.018*"late" + '
  '0.018*"pleasant" + 0.018*"fantastic"'),
 (1,
  '0.048*"bed" + 0.040*"small" + 0.039*"bathroom" + 0.033*"value" + '
  '0.025*"kitchen" + 0.025*"money" + 0.023*"bedroom" + 0.022*"guest" + '
  '0.022*"unit" + 0.017*"respond"'),
 (2,
  '0.105*"time" + 0.062*"amazing" + 0.044*"expect" + 0.041*"book" + '
  '0.035*"wonderful" + 0.026*"friend" + 0.025*"stylish" + 0.025*"new" + '
  '0.020*"condo" + 0.017*"better"'),
 (3,
  '0.069*"get" + 0.041*"go" + 0.038*"day" + 0.033*"find" + 0.032*"use" + '
  '0.032*"airbnb" + 0.029*"good" + 0.022*"want" + 0.022*"could" + '
  '0.021*"thing"'),
 (4,
  '0.195*"apartment" + 0.092*"well" + 0.033*"work" + 0.026*"picture" + '
  '0.024*"building" + 0.022*"awesome" + 0.015*"show" + 0.015*"property" + '
  '0.014*"kid" + 0.012*"communicate"'),
 (5,
  '0.084*"location" + 0.051*"mrt" + 0.031*"easy" + 0.030*"close" + '
  '0.030*"station" + 0.028*"walk" + 0.025*"convenient" + 0.022*"food" + '
  '0.019*"area" + 0.017*"lot"'),
 (6,
  '0.084*"clean" + 0.076*"host" + 0.067*"room" + 0.035*"really" + '
  '0.035*"would" + 0.029*"thank" + 0.027*"comfortable" + 0.027*"need" + '
  '0.024*"everything" + 0.021*"helpful"'),
 (7,
  '0.070*"responsive" + 0.050*"response" + 0.039*"available" + 0.030*"wor'
  'th" + '
  '0.028*"enjoy" + 0.027*"comfy" + 0.025*"able" + 0.023*"cleanliness" + '
  '0.022*"reply" + 0.021*"free"'),
 (8,
  '0.038*"bit" + 0.037*"little" + 0.036*"price" + 0.027*"door" + '
  '0.026*"however" + 0.023*"may" + 0.022*"hotel" + 0.017*"bad" + 0.017*"t'
  'hink' + '
  '+ 0.014*"open"'),
 (9,
  '0.155*"singapore" + 0.060*"make" + 0.050*"home" + 0.042*"feel" + '
  '0.038*"place" + 0.030*"beautiful" + 0.028*"trip" + 0.026*"visit" + '
  '0.020*"hospitality" + 0.020*"welcome"')]
```

```
In [273]: # Compute Perplexity for 10 topics
print('\nPerplexity: ', lda_model.log_perplexity(corpus))
        # a measure of how good the model is. lower the better.

# Compute Coherence Score for 10 topics
coherence_model_lda = CoherenceModel(model=lda_model, texts=data_lemmatized,
coherence_lda = coherence_model_lda.get_coherence())
print('\nCoherence Score: ', coherence_lda)
ne python sum builtin instead.
    score += np.sum(cnt * logsumexp(Elogtheta + Elogbeta[:, int(id)])) for
id, cnt in doc)
/anaconda3/lib/python3.6/site-packages/gensim/models/ldamodel.py:826: Dep
recationWarning: Calling np.sum(generator) is deprecated, and in the futu
re will give a different result. Use np.sum(np.from_iter(generator)) or t
he python sum builtin instead.
    score += np.sum(cnt * logsumexp(Elogtheta + Elogbeta[:, int(id)])) for
id, cnt in doc)
/anaconda3/lib/python3.6/site-packages/gensim/models/ldamodel.py:826: Dep
recationWarning: Calling np.sum(generator) is deprecated, and in the futu
re will give a different result. Use np.sum(np.from_iter(generator)) or t
he python sum builtin instead.
    score += np.sum(cnt * logsumexp(Elogtheta + Elogbeta[:, int(id)])) for
id, cnt in doc)
```

Perplexity: -7.362204337574667

Coherence Score: 0.5647060745372311

Visualizing Topic Modelling

```
In [44]: # Plotting tools
import pyLDAvis
import pyLDAvis.gensim # don't skip this
import matplotlib.pyplot as plt
%matplotlib inline
```

```
In [276]: # Visualize the topics for 10 topics
pyLDAvis.enable_notebook()
vis = pyLDAvis.gensim.prepare(lda_model, corpus, id2word)
vis
```

/anaconda3/lib/python3.6/site-packages/pyLDAvis/_prepare.py:257: FutureWarning: Sorting because non-concatenation axis is not aligned. A future version

of pandas will change to not sort by default.

To accept the future behavior, pass 'sort=False'.

To retain the current behavior and silence the warning, pass 'sort=True'.

```
return pd.concat([default_term_info] + list(topic_dfs))
```

Out[276]:
Typesetting math: 0%

Refine the model by trying different number of topics

Try different number of topics - aim to get the best perplexity and coherence scores, and where the topics mapped out during visualization do not overlap each other and are as far away from each other as possible.

```
In [39]: # Build LDA model for 5 topics
lda_model = gensim.models.ldamodel.LdaModel(corpus=corpus,
                                             id2word=id2word,
                                             num_topics=5,
                                             random_state=100,
                                             update_every=1,
                                             chunksize=100,
                                             passes=10,
                                             alpha='auto',
                                             per_word_topics=True)
```

```
In [40]: # Print the Keyword in the 5 topics

from pprint import pprint

pprint(lda_model.print_topics())
doc_lda = lda_model[corpus]

[(0,
  '0.084*apartment' + 0.039*well' + 0.029*area' + 0.024*perfect' + '
  '0.022*amazing' + 0.018*view' + 0.017*locate' + 0.017*excellent' +
  ,
  '0.017*pool' + 0.016*communication'),
 (1,
  '0.019*get' + 0.017*use' + 0.016*day' + 0.014*night' + 0.012*work'
+ '
  '0.012*could' + 0.011*thing' + 0.010*guest' + 0.010*airbnb' + '
  '0.009*sleep'),
 (2,
  '0.111*room' + 0.035*bed' + 0.029*small' + 0.029*bathroom' + 0.024
*"big" '
  '+ 0.022*space' + 0.021*overall' + 0.020*quite' + 0.018*expect' + '
  '0.018*people'),
 (3,
  '0.098*location' + 0.059*mrt' + 0.036*easy' + 0.035*close' + '
  '0.035*station' + 0.033*walk' + 0.029*convenient' + 0.026*food' + '
  '0.020*lot' + 0.019*bus'),
 (4,
  '0.055*clean' + 0.050*host' + 0.028*singapore' + 0.023*really' + '
  '0.023*check' + 0.023*would' + 0.019*thank' + 0.018*comfortable' +
  ,
  '0.018*need' + 0.017*time')]
```

```
In [41]: # Compute Perplexity for 5 topics
print('\nPerplexity: ', lda_model.log_perplexity(corpus))
        # a measure of how good the model is. lower the better.

# Compute Coherence Score for 5 topics
coherence_model_lda = CoherenceModel(model=lda_model, texts=data_lemmatized,
coherence_lda = coherence_model_lda.get_coherence()
print('\nCoherence Score: ', coherence_lda)
```

Perplexity: -7.210659748379875

Coherence Score: 0.6214173456128512

```
In [45]: # Visualize the topics for 5 topics
pyLDAvis.enable_notebook()
vis = pyLDAvis.gensim.prepare(lda_model, corpus, id2word)
vis
```

/anaconda3/lib/python3.6/site-packages/pyLDAvis/_prepare.py:257: FutureWarning: Sorting because non-concatenation axis is not aligned. A future version of pandas will change to not sort by default.

To accept the future behavior, pass 'sort=False'.

To retain the current behavior and silence the warning, pass 'sort=True'.

```
return pd.concat([default_term_info] + list(topic_dfs))
```

Out[45]:

```
In [281]: # Build LDA model for 4 topics
lda_model = gensim.models.ldamodel.LdaModel(corpus=corpus,
                                             id2word=id2word,
                                             num_topics=4,
                                             random_state=100,
                                             update_every=1,
                                             chunksize=100,
                                             passes=10,
                                             alpha='auto',
                                             per_word_topics=True)
```

The history saving thread hit an unexpected error (OperationalError('unable to open database file',)).History will not be written to the database.

```
In [282]: # Print the Keyword in the 4 topics
```

```
from pprint import pprint
```

```
pprint(lda_model.print_topics())
doc_lda = lda_model[corpus]
```

```
[(0,
  '0.052*"mrt" + 0.031*"close" + 0.031*"station" + 0.029*"walk" + '
  '0.025*"convenient" + 0.023*"food" + 0.018*"area" + 0.017*"bus" + '
  '0.017*"restaurant" + 0.015*"minute"'),
 (1,
  '0.060*"room" + 0.019*"bed" + 0.016*"small" + 0.015*"bathroom" + 0.013
*"use" '
  '+ 0.012*"night" + 0.010*"get" + 0.010*"kitchen" + 0.010*"money" + '
  '0.009*"work"'),
 (2,
  '0.044*"singapore" + 0.035*"really" + 0.029*"thank" + 0.026*"time" + '
  '0.020*"family" + 0.019*"stay" + 0.019*"go" + 0.019*"house" + '
  '0.017*"perfect" + 0.017*"make"'),
 (3,
  '0.061*"location" + 0.058*"clean" + 0.053*"host" + 0.040*"apartment" +
  '0.024*"check" + 0.024*"would" + 0.019*"well" + 0.019*"need" + '
  '0.018*"comfortable" + 0.017*"everything"')]
```

Typesetting math: 0%

```
In [283]: # Compute Perplexity for 4 topics
print('\nPerplexity: ', lda_model.log_perplexity(corpus))
        # a measure of how good the model is. lower the better.

# Compute Coherence Score for 4 topics
coherence_model_lda = CoherenceModel(model=lda_model, texts=data_lemmatized,
coherence_lda = coherence_model_lda.get_coherence()
print('\nCoherence Score: ', coherence_lda)
```

/anaconda3/lib/python3.6/site-packages/gensim/models/ldamodel.py:826: DeprecationWarning: Calling np.sum(generator) is deprecated, and in the future will give a different result. Use np.sum(np.from_iter(generator)) or the python sum builtin instead.

score += np.sum(cnt * logsumexp(Elogthetaad + Elogbeta[:, int(id)])) for id, cnt in doc)

/anaconda3/lib/python3.6/site-packages/gensim/models/ldamodel.py:826: DeprecationWarning: Calling np.sum(generator) is deprecated, and in the future will give a different result. Use np.sum(np.from_iter(generator)) or the python sum builtin instead.

score += np.sum(cnt * logsumexp(Elogthetaad + Elogbeta[:, int(id)])) for id, cnt in doc)

/anaconda3/lib/python3.6/site-packages/gensim/models/ldamodel.py:826: DeprecationWarning: Calling np.sum(generator) is deprecated, and in the future will give a different result. Use np.sum(np.from_iter(generator)) or the python sum builtin instead.

score += np.sum(cnt * logsumexp(Elogthetaad + Elogbeta[:, int(id)])) for id, cnt in doc)

/anaconda3/lib/python3.6/site-packages/gensim/models/ldamodel.py:826: Dep

```
In [284]: # Visualize the topics for 4 topics
pyLDavis.enable_notebook()
vis = pyLDavis.gensim.prepare(lda_model, corpus, id2word)
vis
```

/anaconda3/lib/python3.6/site-packages/pyLDavis/_prepare.py:257: FutureWarning: Sorting because non-concatenation axis is not aligned. A future version of pandas will change to not sort by default.

To accept the future behavior, pass 'sort=False'.

To retain the current behavior and silence the warning, pass 'sort=True'.

```
return pd.concat([default_term_info] + list(topic_dfs))
```

Out[284]: